

```
import numpy as np
import pandas as pd
import matplotlib as plt
%matplotlib inline
import seaborn as sns
```

```
df = pd.read_csv('diabetes.csv')
df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                              768 non-null    int64
2   BloodPressure                        768 non-null    int64
3   SkinThickness                        768 non-null    int64
4   Insulin                              768 non-null    int64
5   BMI                                  768 non-null    float64
6   DiabetesPedigreeFunction              768 non-null    float64
7   Age                                  768 non-null    int64
8   Outcome                              768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
df.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabetes
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.347107
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.477208


```
df.shape
```

```
(768, 9)
```

```
df.Outcome.value_counts()
```

```
0    500
1    268
Name: Outcome, dtype: int64
```

```
sns.countplot(x='Outcome' , data =df);
```



```
df.dtypes
```

```
Pregnancies      int64
Glucose           int64
BloodPressure     int64
SkinThickness     int64
Insulin           int64
BMI               float64
DiabetesPedigreeFunction float64
Age               int64
Outcome           int64
dtype: object
```

```
df.describe()
```

```
#there are incorrect values i.e. 0 in Glucose, BloodPressure, SkinThickness, Insulin, BMI.
# replacing 0 with median of corresponding column.
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabetes
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

```
dataframe_temp = df.drop(["Pregnancies","Outcome"],axis = 1)
dataframe_temp
medians = dataframe_temp.median()
print("medians",medians)
dataframe_nonzero = dataframe_temp.replace(0,medians)
dataframe_nonzero["Pregnancies"] = df["Pregnancies"]
dataframe_nonzero["Outcome"] = df["Outcome"]
dataframe_nonzero
```

```

medians Glucose      117.0000
BloodPressure      72.0000
SkinThickness      23.0000
Insulin            30.5000
BMI                32.0000
DiabetesPedigreeFunction  0.3725
Age                33.0000

```

```

corr = dataframe_nonzero.corr()
corr

```

	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Pregnancies	Outcome
Glucose	1.000000	0.218937	0.172143	0.357573	0.231400	0.137327	0.266909	0.128213	0.492782
BloodPressure	0.218937	1.000000	0.147809	-0.028721	0.281132	-0.002378	0.324915	0.208615	0.165723
SkinThickness	0.172143	0.147809	1.000000	0.238188	0.546951	0.142977	0.054514	0.032568	0.189065
Insulin	0.357573	-0.028721	0.238188	1.000000	0.189022	0.178029	-0.015413	-0.055697	0.148457
BMI	0.231400	0.281132	0.546951	0.189022	1.000000	0.153506	0.025744	0.021546	0.312249
DiabetesPedigreeFunction	0.137327	-0.002378	0.142977	0.178029	0.153506	1.000000	0.266909	0.128213	0.492782
Age	0.266909	0.324915	0.054514	-0.015413	0.025744	0.266909	1.000000	0.128213	0.492782
Pregnancies	0.128213	0.208615	0.032568	-0.055697	0.021546	0.128213	0.128213	1.000000	0.492782
Outcome	0.492782	0.165723	0.189065	0.148457	0.312249	0.492782	0.492782	0.492782	1.000000

Takeaway : outcome is positively corelated to Glucose feature.


Age & no. of pregencies have positive corelation.

BMI & Skin thickness has positive corelation

No other strong negative corelation is observed.

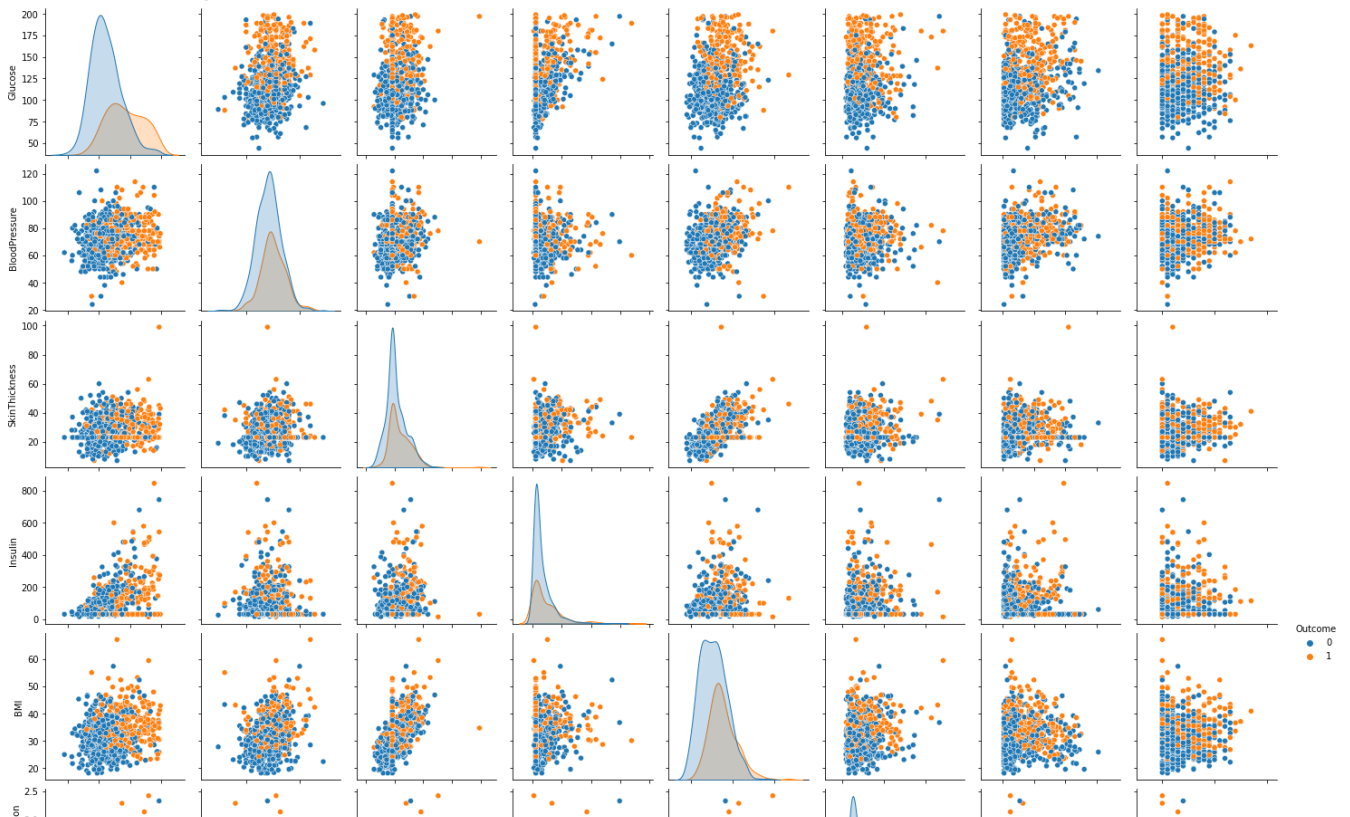
```
sns.heatmap(corr)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f945b12a150>
```



```
sns.pairplot(dataframe_nonzero, diag_kind='kde', hue="Outcome") # plotting pairplot
```

<seaborn.axisgrid.PairGrid at 0x7f9458409390>



```
from sklearn.model_selection import train_test_split
X = dataframe_nonzero.drop('Outcome', axis=1)
Y = dataframe_nonzero['Outcome']
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.30, random_state=1)
```



Training Support vector Machines



```
from sklearn import svm
from sklearn.svm import SVC

clf = svm.SVC(C = 100, gamma= "scale")
clf.fit(X_train,Y_train)
```

SVC(C=100)

```
score1 = clf.score(X_test,Y_test)
score1
```

0.7705627705627706

```
from sklearn import metrics
Y_pred = clf.predict(X_test)
print( metrics.confusion_matrix(Y_test,Y_pred))
```

[[131 15]

```
[ 38  47]]
```

Scaling the datapoints using MinMax Scalar

```
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.fit_transform(X_test)

#Zscore
from scipy.stats import zscore
X_train_z = X_train.apply(zscore) # converting to Z score
X_test_z = X_test.apply(zscore)

# Model score on Minmax scaled values
clf = svm.SVC(C = 10,gamma= "scale")
clf.fit(X_train_scaled,Y_train)
score2 = clf.score(X_test_scaled,Y_test)
score2
```

```
0.7532467532467533
```

```
# Model score using zscore values
clf = svm.SVC(C = 10,gamma= "scale")
clf.fit(X_train_z,Y_train)
score3 = clf.score(X_test_z,Y_test)
score3
```

```
0.7272727272727273
```

We can try increasing either C or gamma to fit a more complex model.

```
clf = svm.SVC(C = 1000,gamma= "scale")
clf.fit(X_train,Y_train)
score4 = clf.score(X_test,Y_test)
print("Model score for non-scaled datapoints", score4)
```

model accuracy has increased on non-scaled data,however for scaled values with c = 1000, mo

```
Model score for non-scaled datapoints 0.7835497835497836
```

```
import multiprocessing
from sklearn.model_selection import GridSearchCV
```

```
param_grid = [
    {
        'kernel': ['linear', 'rbf', 'poly', 'sigmoid'],
```

```
'C': [ 0.1, 0.2, 0.4, 0.5, 1.0, 1.5, 1.8, 2.0, 2.5, 3.0 ]    } ]
```

```
gs = GridSearchCV(estimator=SVC(), param_grid=param_grid,scoring='accuracy', cv=10, n_jobs=mu
```

```
gs.fit(X_train_scaled, Y_train)
```

```
GridSearchCV(cv=10, estimator=SVC(), n_jobs=2,
             param_grid=[{'C': [0.1, 0.2, 0.4, 0.5, 1.0, 1.5, 1.8, 2.0, 2.5,
                                3.0],
                           'kernel': ['linear', 'rbf', 'poly', 'sigmoid']}],
             scoring='accuracy')
```

```
gs.best_estimator_
```

```
SVC(kernel='linear')
```

```
gs.best_score_
```

```
0.7653738644304683
```

Calculate AUC score and plot ROC curve

```
from sklearn.metrics import roc_auc_score,roc_curve
```

```
auc = roc_auc_score(Y_test,Y_pred)
print("AUC %0.3f" %auc)
```

```
AUC 0.725
```

AUC measures how true positive rate (recall) and false positive rate trade off

```
from sklearn import model_selection
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
```

```
seed = 7
models = []
models.append(('LR', LogisticRegression()))
models.append(('LDA', LinearDiscriminantAnalysis()))
models.append(('KNN', KNeighborsClassifier()))
models.append(('CART', DecisionTreeClassifier()))
```



```
models.append(('NB', GaussianNB()))
models.append(('SVM', SVC()))
models.append(('RFC', RandomForestClassifier()))
results = []
names = []
scoring = 'accuracy'
import warnings
warnings.filterwarnings("ignore")
for name, model in models:
    kfold = model_selection.KFold(n_splits=10, random_state=None)
    cv_results = model_selection.cross_val_score(model, X_train, Y_train,
cv=kfold, scoring=scoring)
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)

LR: 0.763347 (0.061653)
LDA: 0.770790 (0.062164)
KNN: 0.696261 (0.085826)
CART: 0.683368 (0.060197)
NB: 0.730119 (0.054511)
SVM: 0.744794 (0.054232)
RFC: 0.757757 (0.055708)
```

Linear Discriminant Analysis is giving a better accuracy of 77% as compared with the other models.